

Combinatorial Mathematics

Mong-Jen Kao (高孟駿)

Monday 18:30 – 21:20

Outline

- The Maximum Matching Problem
 - A Generic Algorithm and the Berge's Theorem
 - Solving the Augmenting Path Problem
 - DFS-based & BFS-based Algorithms for Bipartite Graphs
 - The Blossom Algorithm for General Graphs
- Concluding Notes
 - The best algorithms for Maximum Matching

Characterization of Bipartite Graphs

Identify the two partite sets of a bipartite graph when it is not given.

Characterization of Bipartite Graphs

- The following theorem is simple and intuitive to prove.

Theorem. (Characterization of Bipartite Graphs)

A graph $G = (V, E)$ is bipartite if and only if it has a 2-coloring, i.e., a 2-coloring for V such that no edge $e \in E$ is monochromatic.

- Note that, the 2-colorability of G can be tested by a simple DFS.
 - If G has a 2-coloring, then it also corresponds to a valid classification of the two partite sets.

You will need this fact in ProgHW #1.

An Alternative BFS-based Algorithm

An Alternative Algorithm

- Let X_0 be the set of all unmatched vertices in G .
- For any $i = 0, 1, 2, \dots$, define
 - X_{2i+1} to be the set of *unvisited* vertices (not in $X_{\leq 2i}$) that can be reached from X_{2i} **using an edge not in M** .
 - X_{2i+2} to be the set of *unvisited* vertices (not in $X_{\leq 2i+1}$) that can be reached from X_{2i+1} **using an edge in M** .

An Alternative Algorithm


- Let X_0 be the set of all unmatched vertices in G .
- Formally, for any $i = 0, 1, 2, \dots$, define

$$X_{2i+1} := \{ v \in V \setminus X_{\leq 2i} : \exists u \in X_{2i} \text{ s.t. } (u, v) \notin M \}$$

and

$$X_{2i+2} := \{ v \in V \setminus X_{\leq 2i+1} : \exists u \in X_{2i+1} \text{ s.t. } (u, v) \in M \}.$$


The *Alternating Forest* Formed by X_i



The *Alternating Forest* Formed by X_i


- The roots are the unmatched vertices in X_0 .
 - Each non-root vertex has exactly one parent in the forest.
- For any vertex $v \in V$,
 - Let $P^\uparrow(v)$ be the path from v to its root in the forest.
 - Also, let $P^\downarrow(v)$ be the path from its root to v in the forest.
- Note that, $P^\uparrow(v)$ and $P^\downarrow(v)$ are uniquely defined, and they are M -alternating paths.

The *Alternating Forest* Formed by X_i



The *Alternating Forest*

Note that, there may still be edges between a vertex in the odd level and other vertices, but we don't care.



- Let $G = (V, E)$ be a bipartite graph and M be a matching for G .

Another BFS-based Augmenting Path Algorithm (for Bipartite Graphs).

1. Let X_0 be the set of unmatched vertices and $t \leftarrow 0$.
2. Repeat until $X_{\leq 2t} = V$, do
 - If there exists an edge $(u, v) \in E$ for some $u, v \in X_{2t}$,
then return the path $\{P^\downarrow(u), P^\uparrow(v)\}$.
 - Otherwise,
form X_{2t+1} and X_{2t+2} as described and set $t \leftarrow t + 1$.
3. Report $\bigcup_{i \geq 0} X_{2i+1}$ as a *vertex cover* with size $|M|$.

The Augmenting Path Problem

in General Graphs

For general graphs,
the augmenting path problem can be solved in $O(nm)$ time via proper vertex contractions.


The Augmenting Path Problem in General Graphs

- Let $G = (V, E)$ be a general graph and M be a matching for G .
- We introduce an algorithm that computes in $O(nm)$ time either
 - An M -augmenting path for G , or,
 - A structure (proof) showing that M is maximum.
Hence, no M -augmenting path exists in the graph.

Note that, we can no longer count on vertex covers for this, since the **strong duality does not hold** between matchings and vertex covers in general graphs.


Blossom, Stem, and Flowers

- A **blossom** is a cycle C with an odd length and $\lfloor |C|/2 \rfloor$ matched edges in M .
 - The vertex $v \in C$ that is not incident to any matched edge is called the “**joint**” of the blossom.




Blossom, Stem, and Flowers

- A stem is an M -alternating path
with an even length and ends at a matched edge in M .





Blossom, Stem, and Flowers

- A flower is a stem and a blossom
such that the stem ends at the joint of the blossom.



Contracting a Blossom


- Let C be a blossom in G .
 - Define G_C to be the graph obtained by contracting C in G , and M'_C be the remaining set of matched edges.



- Let C be a blossom in G .
 - Define G_C to be the graph obtained by contracting C in G , and M'_C be the remaining set of matched edges.

Lemma. (Blossom Contraction)

G has an M -augmenting path
if and only if G_C has an M'_C -augmenting path.




The Blossom Algorithm (by Jack Edmonds)

- Let X_0 be the set of all unmatched vertices in G .
- For any $i = 0, 1, 2, \dots$, define
 - X_{2i+1} to be the set of *unvisited* vertices (not in $X_{\leq 2i}$) that can be reached from X_{2i} **using an edge not in M** .
 - X_{2i+2} to be the set of *unvisited* vertices (not in $X_{\leq 2i+1}$) that can be reached from X_{2i+1} **using an edge in M** .

The Blossom Algorithm (by Jack Edmonds)

- Consider the alternating forest formed by X_i for all $i \geq 0$.
- If there exists **a cross edge in an even level**,
i.e., $(u, v) \in E$ for some $u, v \in X_{2i}$ and some $i \geq 0$,
then $\{P^\downarrow(u), P^\uparrow(v)\}$ is either *an M-augmenting path* or a *flower*!
 - If $P^\downarrow(u) \cap P^\uparrow(v) = \emptyset$, then it is an augmenting path.
 - Otherwise,
it is a flower with the common part being the stem.

The *Alternating Forest* Formed by X_i



- Let $G = (V, E)$ be a graph and M be a matching for G .

The Blossom Algorithm (by Jack Edmonds).

1. Let X_0 be the set of unmatched vertices and $t \leftarrow 0$.
2. Repeat until $X_{\leq 2t} = V$, do
 - If there exists an edge $(u, v) \in E$ for some $u, v \in X_{2t}$,
 - If $P^\downarrow(u) \cap P^\uparrow(v) = \emptyset$, then return the path $\{P^\downarrow(u), P^\uparrow(v)\}$.
 - Otherwise, let $C \leftarrow P^\downarrow(u) \Delta P^\uparrow(v)$. Apply the algorithm recursively on G_C and M'_C . Expand the result and return it.
 - Otherwise,
form X_{2t+1} and X_{2t+2} as described and set $t \leftarrow t + 1$.
3. Report $\bigcup_{i \geq 0} X_{2i+1}$ as a proof.

The Correctness of the Blossom Algorithm

Analysis of the Algorithm

- For the correctness of the algorithm,
 - It is clear that, when the blossom algorithm returns an M -augmenting path, it is indeed a valid one.
 - We need to show that, when the algorithm returns a proof (reports “No”), M is indeed a maximum matching.

For this, we will use the Tutte-Berge Max-Min Theorem.

Lemma. (Tutte-Berge Max-Min Theorem)

Let $G = (V, E)$ be a graph,

$U \subseteq V$ be a vertex subset, and $M \subseteq E$ be a matching.

Then we always have

$$|M| \leq \frac{|V| + |U| - \text{odd}(G \setminus U)}{2},$$

where $\text{odd}(G \setminus U)$ is the number of components with an odd size in $G \setminus U$.

- Later we will see that, the inequality holds with equality for properly chosen M and U when G contains no blossom.

Let $G = (V, E)$ be a graph, $U \subseteq V$ be a vertex subset, and $M \subseteq E$ be a matching. Then we always have

$$|M| \leq \frac{|V| + |U| - \text{odd}(G \setminus U)}{2},$$

where $\text{odd}(G \setminus U)$ is the number of odd components in $G \setminus U$.

- Consider the components in U and $G \setminus U$.
 - Each vertex in U is incident with at most one edge in M .
 - For the remaining components K in $G \setminus U$,
it contains at most $\left\lfloor \frac{|K|}{2} \right\rfloor$ edges in M .

Since all endpoints of
the edges in M are distinct.

We always have $|M| \leq (|V| + |U| - \text{odd}(G \setminus U))/2$,

where $\text{odd}(G \setminus U)$ is the number of odd components in $G \setminus U$.

■ Consider the components in U and $G \setminus U$.

– Each vertex in U is incident with at most one edge in M .

– For the remaining components K in $G \setminus U$,

it contains at most $\left\lfloor \frac{|K|}{2} \right\rfloor$ edges in M .

Since all endpoints of the edges in M are distinct.

■ Hence,


$$|M| \leq |U| + \sum_i \left\lfloor \frac{|K_i|}{2} \right\rfloor = |U| + \frac{|V| - |U|}{2} - \frac{\text{odd}(G \setminus U)}{2}.$$

Analysis of the Algorithm

- Suppose that the Blossom algorithm returns a proof $U_{i \geq 0} X_{2i+1}$.
- By the Tutte-Berge's inequality, to prove that M is a maximum matching for G , it suffices to show that

The choice of $U := U_{i \geq 0} X_{2i+1}$ will make
the Tutte-Berge's inequality hold with equality.

The *Alternating Forest* Formed by X_i



There is no cross edge in the even levels.

Hence, the vertices in the even levels become **isolated**.

The remaining components (not included in the forest) are perfectly matched by M .

Analysis of the Algorithm

- Let $U := \bigcup_{i \geq 0} X_{2i+1}$.

- Then,

$$\begin{aligned}|M| &= |U| + \frac{|V \setminus X_{\geq 0}|}{2} = \frac{|V| + |U| - |\bigcup_{i \geq 0} X_{2i}|}{2} \\&= \frac{|V| + |U| - \text{odd}(G \setminus U)}{2}.\end{aligned}$$

- Hence, M is a maximum matching for G .

Concluding Notes

Best Algorithm for the Maximum Bipartite Matching

- In this lecture,
we have seen an $O(nm) = O(n^3)$ algorithm for this problem.
- The best algorithm for this problem is the Hopcroft-Karp
algorithm, which runs in $O(\sqrt{nm}) = O(n^{2.5})$.

The Hopcroft-Karp Algorithm

- The idea is to perform a **BFS** simultaneously from all unmatched vertices in one partite set **to form alternating layers** until some unmatched vertices in the other partite set is met.
- Then a **layer-guided DFS** is used to construct a maximal set of vertex-disjoint shortest augmenting paths.
- It is guaranteed that, only $O(\sqrt{n})$ rounds are needed before the maximum matching is computed.

Best Algorithm for the Maximum Bipartite Matching

- This problem is a special case of the max-flow problem.

A number of flow algorithms are applicable.

- Practically,
the most efficient one is the Dinic's algorithm.
 - Theoretically, the best algorithm is the “Almost linear-time”
max-flow algorithm that runs in $m^{1+o(1)}$ time.

Maximum Matching in General Graphs

- For general graphs, we have seen the Edmonds Blossom Algorithm, which runs in $O(n^2m) = O(n^4)$ time.
- The best (and more complicated) algorithm, due to Micali and Vazirani, solves this problem in $O(\sqrt{nm}) = o(n^{2.5})$ time.